

Embedded Spaces for Discovering New Classes

P. C. Harris

Massachusetts Institute of Technology

pcharris@mit.edu

Abstract

A representation can encode a novel class perfectly and still be useless for finding it. On a supervised contrastive encoder trained with one class held out, a linear probe separates that class from the rest at an AUC of 0.94, while the same space, asked to flag those examples without being told what to look for, flags none of them. Representation quality is routinely measured on the first side of this gap — linear probes, k -NN, transfer — and scientific discovery lives entirely on the second. We study the standard open-world loop (score a corpus against a labelled reference, pool the most anomalous points, cluster the pool into candidate class anchors, adapt the representation) and take as the primary metric the fraction of that pool which is genuinely novel, since every later stage is downstream of it. We contribute four things. First, an algebra for the pool: purity is $\Pi = \varepsilon_s b/q$ with a ceiling $\min(1, b/q)$ that depends on the pooling threshold alone, which separates “the representation is bad” from “the threshold is bad” — questions the literature does not currently distinguish. Second, a label-free rule for choosing that threshold, derived from the same Neyman–Pearson density ratio that scores the corpus; it recovers the oracle threshold exactly on synthetic data, and where the scorer is too weak to support a choice it declines to pool and reports the refusal as an output. Third, evidence that ranking training objectives by discoverability is not the same as ranking them by probe accuracy: adding a sliced-isotropy regulariser to a supervised contrastive objective is a tie on the linear probe and costs 1.1 points of top-1 accuracy, while raising novelty-detection AUC from 0.233 to 0.620. Fourth, a 2×2 decomposition of representation construction against the discovery loop, which finds the two to be substitutes more often than complements. Throughout we separate the $\sim 1\%$ and $\sim 10\%$ novelty regimes, which invert several conclusions; we report where the pipeline fails as carefully as where it works; and we rebuild every CIFAR result from random initialisation, because a hub-pretrained backbone has already seen the class it is being asked to discover.

1 Introduction

A representation can encode a novel class perfectly and still be useless for finding it. This is not a hypothetical failure mode; it is the ordinary situation. Train a supervised contrastive encoder on a set of known classes, hold one class out, and a linear probe trained to separate the held-out class from the rest reaches an AUC of 0.94. Now ask the same space a different question: given a *single* image and no knowledge of what is being looked for, is this one novel? Run as a hypothesis test at a 5% false-positive rate, that space flags none of them — its detection rate on genuinely novel images is no better than the false-positive rate itself. The information is present, and linearly accessible. It is simply not accessible to anyone who does not already have the answer.

The gap between those two numbers is the subject of this paper. It is the gap between a representation being *decodable* and being *discoverable*, and it matters because scientific discovery lives entirely on the second side of it. A search for a new particle, a new galaxy morphology, or a new failure mode of a deployed model does not come with labels for the thing being searched for. What it has is a large corpus, a smaller reference sample of what is already understood, and a question: is there something here that the reference does not explain?

Most of the representation-learning literature optimises for the first side. Linear probe accuracy, k -NN accuracy, and transfer benchmarks all measure how well a space supports a decision that someone else has already specified. They are the right metrics for that purpose. They are close to uninformative about whether a space will surface structure nobody has named yet, and — as the

two numbers above indicate — they can be actively misleading, because the training objectives that maximise them are not the objectives that make distances mean anything in absolute terms.

What we do. We treat discovery as the target and work backwards. The pipeline we study is the standard open-world loop: score every point in a corpus against a reference of labelled examples, take the most anomalous ones as a *pool*, cluster the pool to propose new class anchors, and adapt the representation to the proposal. Every stage of that loop is downstream of one quantity — the fraction of the pool that is genuinely novel, which we call *pool purity*. Purity is the only forward-looking metric in the pipeline: every other measure describes a space you already have, whereas purity predicts whether the next step will improve that space or corrupt it. When purity is low the loop does not merely fail to help. It clusters background, plants anchors on that background, and then fine-tunes the representation to make the spurious structure sharper — so a bad pool leaves the representation worse than not running the loop at all.

Contributions.

1. **An algebra for the pool, and a ceiling.** Writing b for the novel base rate, q for the pool fraction, and ε_s for the fraction of novel points retained, pool purity is $\Pi = \varepsilon_s b / q$, so the enrichment Π / b is exactly ε_s / q and $\Pi \leq \min(1, b / q)$. The ceiling is a property of the cut alone. It says that a pool threshold inherited without derivation — as ours was — can cap achievable purity at a level no representation can rescue, and it separates “the space is bad” from “the cut is bad” for the first time.
2. **A label-free rule for the cut.** Because purity requires the novel labels, tuning the cut on purity is oracle tuning. We derive the novel-count estimator that the Neyman–Pearson critic already provides, and select the tightest cut that retains enough *estimated* novel points to remain clusterable. The rule uses only the corpus and the labelled reference. On synthetic data with an exact critic it recovers the oracle cut exactly; on real embeddings it raises pool purity from 0.232 to 0.778 where the scorer is strong, and refuses — reporting that it has refused — where it is not.
3. **A regime split, stated as a condition rather than a caveat.** Single-class novelty at a rate of $\sim 1\%$ and multi-class novelty at $\sim 10\%$ are different problems with different answers: which pooling statistic wins, whether stricter thresholds help or hurt, and what purity is attainable at all invert between them. We report them separately throughout.
4. **A reduced set of objectives.** From a design cube of ten contrastive objectives we identify the few that earn a place for discovery specifically, and show that the ranking by discovery is not the ranking by probe. On a from-scratch CIFAR-100 lineage, adding sliced-isotropy to a supervised contrastive objective is a tie on the linear probe (0.937 vs. 0.940) and costs 1.1 points of top-1, while raising every detection column — Euclidean AUC $0.233 \rightarrow 0.620$, tied-Mahalanobis $0.279 \rightarrow 0.602$. One point of accuracy buys a $2.7\times$ improvement in the geometry discovery actually runs on.
5. **Two constructions, and a decomposition of what they buy.** We study a residual decomposition — train a second encoder on what the anchors fail to explain, then concatenate — and frozen-space discovery, in which only the anchors move. Because both can be composed with the discovery loop, we report them in a 2×2 ablation (construction $\times$ discovery) rather than as a single delta, which separates a genuine interaction from an additive stack. The composition is superadditive in 12 of 20 cells but concentrated: the effect that takes per-event detection from nominal to $6\times$ nominal lives in one cell, while the two largest gains elsewhere turn out to be main effects of the construction and of discovery separately (Section 4). We state this narrowly because the decomposition is what makes it checkable.
6. **A leakage-free lineage for the claims that need one.** Every CIFAR cell built from hub-pretrained weights inherits an encoder that already saw the held-out class, which makes it unusable as evidence for discovery. We rebuild the CIFAR battery from random initialisation

and report the pretrained cells only as an ablation. Dataset ordering throughout is by distance from pretraining rather than alphabetically, so the reader can see which claims rest on a backbone that could have seen the answer.

What this paper is not. We do not claim that a regulariser which makes embedded distance behave like a log-likelihood thereby improves discovery; we tested that directly and it does not hold (Section 6). We do not claim the pipeline works everywhere. A substantial part of what follows is a characterisation of where it does not, and why — in several cases the obstruction is the novel-class rate rather than anything about the representation, which is a limit no encoder can move.

2 Setup: discovery as a two-sample problem

The open-world protocol. We are given a corpus $\mathcal{D}$ of N embedded points, of which an unknown fraction b belong to classes never seen in training, and a reference $\mathcal{R} \subset \mathcal{D}$ of points carrying known class labels. No label for any novel class is available at any point — not for choosing hyperparameters, not for setting thresholds. We state this explicitly because the constraint is easy to violate accidentally, and because a method that quietly violates it reports numbers that cannot be reproduced in the situation it claims to address.

The loop. One round proceeds as follows. Score every point by an anomaly score $s(x)$; threshold to obtain a pool $\mathcal{P}$; cluster $\mathcal{P}$ and append the resulting centroids to the anchor set; pseudo-label every pooled point by its nearest discovered anchor; and adapt the representation (or, in the frozen variant of Section 4, only the anchors) to the enlarged anchor set. The round repeats with the updated anchors.

What we measure, and what the numbers mean. Because this paper compares two kinds of quality that are usually not distinguished, it is worth stating plainly what each number is before any of them appear.

Decodability measures how well a space supports a decision somebody has already specified. We report **linear probe AUC** (train a linear classifier to separate the held-out class from the rest, *using the held-out labels*, and measure how well it does) and **top-1 accuracy** on the seen classes. A probe is an oracle measurement: it is allowed the answer. It says whether the information about the novel class survives in the representation at all, which is a genuine and useful thing to know — and is not what a discovery pipeline can use.

Discoverability measures what the space supports *without* the answer. **Euclidean AUC** ranks every corpus point by its distance to the nearest known-class anchor and asks how well that ranking separates novel from seen; **tied Mahalanobis** does the same with a shared covariance rather than raw distance. Both are ranking measures: they say a novel point tends to score higher than a seen one, but not by how much or where to put a threshold. **Per-event power** is the strictest question we ask. Fix a decision threshold that admits $\alpha = 5\%$ of *reference* points, then report the fraction of genuinely novel points that exceed it. This is the power of a hypothesis test at size α , applied one event at a time. It is bounded below by α in the sense that a space carrying no usable signal scores $\approx \alpha$; a value at or under 0.05 therefore means the space cannot flag an individual novel example at all, and the interesting question is whether a method crosses that floor rather than by how much it improves. Finally, **pool purity** Π is defined below and is the metric this paper treats as primary.

Pool purity, and why everything is downstream of it. Let Π denote the fraction of $\mathcal{P}$ that is truly novel. The pool feeds directly into a clustering step with no supervision downstream of it, so $\mathcal{P}$ is the last point at which a mistake can still be corrected and nothing after it checks the work. At $\Pi \approx 0.01$ the clustering is not finding the novel class; it is partitioning background, and the subsequent fine-tune sharpens that partition. Empirically the probe drops by 0.031 to 0.038 wherever such pools form. We therefore report Π as the primary metric and compare spaces and procedures by it directly. We deliberately do *not* impose a purity threshold: purity is monotone in the cut (Section 3.3), so any fixed level is a statement about the operating point rather than about the space, and no experiment in this paper branches on one. Where a purity is low enough that the subsequent clustering is partitioning background, we say so and give the number.

3 Algorithms

This section describes the objectives that produce the embedding, the statistic that scores it, and the rule that turns that statistic into a pool. The three are usually treated as separate literatures; a contribution of this work is that the middle one — the Neyman–Pearson density ratio — is the same object in all three roles, which is what makes the last one derivable rather than tuned.

3.1 Objectives

We work within a design cube spanning $\{\text{instance, supervised}\} \times \{\text{cosine, bilinear, distance}\} \times \{\text{critics} \times \{\text{softmax, NPLM}\} \times \{\text{estimators} \times \{\text{none, SIGReg, class-wise SIGReg}\} \times \{\text{marginals, which contains SimCLR, SupCon, LeJEPa/VISReg and the density-ratio objectives as corners. Of the ten arms we ran, three earn a place for discovery.}$

SupCon + SIGReg (strong). Supervised positives, cosine critic, softmax estimator, and a sliced isotropic-Gaussian regulariser at weight $\lambda = 5$. This is the discovery workhorse: it produces the highest pool purities we measure. Two points of care. First, λ matters and is not a formality — the same objective at the library default $\lambda = 1$ is a different algorithm and its numbers must not be pooled with these (Appendix B). Second, we make no claim that SIGReg helps *because* it renders distance a log-likelihood; that mechanism is tested and rejected in Section 6. SIGReg earns its place empirically, and structurally: under an isotropic marginal Euclidean distance is Mahalanobis distance, which is what licenses the unit-variance clustering criterion used downstream.

Supervised distance-NPLM. Supervised positives, distance critic, and the NPLM estimator

$$\mathcal{L}[g] = \mathbb{E}_{\mathcal{R}}[e^g - 1] - \mathbb{E}_+[g], \quad (1)$$

whose minimiser is the *absolute* log density ratio, fixed by the calibration $\mathbb{E}_{\mathcal{R}}[e^{g^*}] = 1$ with no free additive constant. Unlike a softmax-normalised critic, this makes a single embedded distance a statement about relative likelihood rather than a statement only about rank. It is the strongest arm in the low-class-count regime and the most robust across pretraining backbones.

SupCon (plain). Supervised positives, cosine critic, softmax estimator, no marginal. Retained as the baseline and as the parent for the residual construction. It is the best space in the paper by linear probe and by semantic agreement, and one of the weakest by every calibration measure — which is the dissociation this paper is organised around, in a single arm.

We drop plain SimCLR, LeJEPa/SIGReg-SSL, instance NPLM-bilinear and instance NPLM-distance to a single ablation row: none reaches the Pareto front for discovery, and the bilinear critic in particular carries roughly three times the seed variance of the distance critic.

3.2 The Neyman–Pearson statistic, in three roles

Let $\mathcal{R}$ be the reference and $\mathcal{D}$ the corpus, with $w = N_{\mathcal{D}}/N_{\mathcal{R}}$. The empirical NP objective

$$\mathcal{L}[f] = w \sum_{x \in \mathcal{R}} (e^{f(x)} - 1) - \sum_{x \in \mathcal{D}} f(x) \quad (2)$$

has minimiser $f^*(x) = \log(p_{\mathcal{D}}(x)/p_{\mathcal{R}}(x))$, obtained by setting the functional derivative $p_{\mathcal{R}}e^f - p_{\mathcal{D}}$ to zero, and therefore satisfies

$$\mathbb{E}_{\mathcal{R}}[e^{f^*}] = \int p_{\mathcal{R}} \cdot \frac{p_{\mathcal{D}}}{p_{\mathcal{R}}} = 1. \quad (3)$$

Equation 2 is Equation 1 with pairs replaced by events. The same functional appears as a contrastive *loss* over pairs, as a two-sample *test statistic* $t_{NP} = -2\mathcal{L}$ distributed as a Wilks log-likelihood ratio, and — below — as a *pool scorer*. We parameterise f as a small kernel ensemble with learnable centres and amplitudes and anneal the bandwidth from the median pairwise distance downwards; a fixed bandwidth is calibrated only for one embedding scale and goes blind when the geometry is rescaled.

Equation 3 is also a free diagnostic, and we report it. A fitted critic need not satisfy it, and a value far from 1 means the numbers built from that fit are unreliable. Measured on trained spaces it is 0.997 to 1.019 in-sample; out-of-sample it ranges to 2.7, reflecting mild overfitting of the reference subsample.

Why the ratio and not the distance. Distance-based pooling asks which points are far from everything. The density ratio asks which points are over-represented relative to what the reference explains. These come apart, and the difference is decisive: a novel class is not a scatter of far-flung outliers but a *dense clump* sitting where the reference has no mass. On a 100-class benchmark the extreme distance tail is owned by background outliers, and tightening the distance threshold makes purity *worse* ($0.036 \rightarrow 0.002$), whereas tightening the ratio threshold makes it better ($0.232 \rightarrow 0.358$). Same space, same rate, opposite response.

3.3 Choosing the pool: an algebra and a rule

The algebra. Let $q = |\mathcal{P}|/N$ and let $\varepsilon_s, \varepsilon_b$ be the fractions of novel and background points retained. Then $q = \varepsilon_s b + \varepsilon_b(1 - b)$ and

$$\Pi = \frac{\varepsilon_s b}{q}, \quad E \equiv \frac{\Pi}{b} = \frac{\varepsilon_s}{q}, \quad \Pi \leq \min\left(1, \frac{b}{q}\right). \quad (4)$$

Three consequences. The enrichment is exactly the lift at operating point q . Purity increases monotonically as q shrinks, so there is no interior optimum and targeting a purity *level* is the wrong objective. And the ceiling $\min(1, b/q)$ is a property of the cut alone: at $q = 0.05$ and $b = 0.01$ no scorer, however good, can exceed $\Pi = 0.20$. A threshold chosen without reference to b therefore caps purity independently of the representation — and the widely used “95th percentile of the reference” is such a threshold.

The rule. Purity cannot be used to choose q , since computing it requires the novel labels. The critic supplies the missing quantity. Under the contamination model $p_{\mathcal{D}} = (1 - b)p_{\mathcal{R}} + bp_S$, with $r = e^f$, the posterior that a corpus point is novel is $1 - (1 - b)/r$; dropping the unknown b and taking the positive part gives the label-free weight

$$w(x) = \left[1 - \frac{1}{r(x)}\right]_+, \quad \mathbb{E}_{\mathcal{D}}[w] = \text{TV}(p_{\mathcal{D}}, p_{\mathcal{R}}) = b \text{TV}(p_S, p_{\mathcal{R}}) \leq b, \quad (5)$$

so $\sum_{x \in S} w(x)$ estimates how many novel points a subset S contains, exactly when the novel class is disjoint from the reference and biased *downward* under overlap — a safe direction, since it tightens the cut. Since purity rises as q falls, the binding constraint is not purity but whether a cluster survives, and that is what we target: ordering the corpus by f and writing $\hat{N}(k)$ for the null-corrected sum of w over the top k points,

$$k^* = \min \{k : \hat{N}(k) \geq n_{\min}\}, \quad q^* = k^*/N, \quad k_{\max} = \left\lfloor \hat{N}(N) / n_{\min} \right\rfloor. \quad (6)$$

The same weights bound the clustering search range, replacing a bound that in the original pipeline was set from the number of held-out classes — a quantity an open-world method does not have.

The null correction, which is not optional. Equation 5 is a positive part, so *any* spread in a fitted f manufactures apparent novelty. On pure-noise scores with no novel class at all, the raw sum returns 665 “novel” points out of 8000 — enough for the rule to fire on nothing and hand the loop a pool of pure background. Subtracting the mean weight of the reference does not repair this, because top-ranked null points carry far above-average weight and their tail survives subtraction. We instead subtract a *rank-matched* null: under the null hypothesis the top q of the corpus is distributed as the top q of the reference, so $\hat{N}(k) = \sum_{i \leq k} w(x_{(i)}) - k \cdot \bar{w}_{\mathcal{R}}^{(q)}$, with $\bar{w}_{\mathcal{R}}^{(q)}$ the mean weight of the top- q fraction of the reference. This cancels to zero under the null while preserving a real excess. When $\hat{N}(N) < n_{\min}$ the rule declines to pool and reports that it has declined; that refusal is an output of the method, and we report it as one.

Constants. $n_{\min}$ and the clipping range on q are set once, globally, from a sweep on held-out embeddings rather than inherited: purity falls monotonically as $n_{\min}$ grows, so it is set as small as the clustering will bear.

4 Constructions

Both constructions in this section are defined *relative to an anchor set*, which is what lets them compose with the discovery loop: discovery enlarges and corrects the anchors, so a construction defined against them is a different construction after discovery than before. That composition is the part of this work with no counterpart we are aware of, and it is also the part most easily overstated, so we measure it as an interaction rather than a delta.

4.1 The residual decomposition

Let μ_y be the anchor for class y and z an embedding. The residual $r = z - \mu_y$ is what the anchor set fails to explain about z . We train a second encoder on r under the same objective family as the parent and either use it alone (*residual*) or concatenate it with the parent (*concat*).

The motivation is that a supervised objective is free to discard any variation that does not separate the classes it was given, and novelty is precisely variation of that kind. The residual target makes that discarded structure the only thing the child can see.

On its own, the residual buys detection at no accuracy cost. On the from-scratch CIFAR-100 lineage, concatenating a residual onto a supervised contrastive parent moves Euclidean novelty AUC from 0.233 to 0.521 while the linear probe moves by +0.0015, a gain two orders of magnitude smaller than the detection gain and of the same order as the probe’s own seed spread ($\sigma = 0.0008$). This is the pattern across cells: the pure residual usually *loses* probe accuracy (four of six strong-parent cells, -0.037 to -0.077), and the concatenation recovers it. We therefore report concat as the construction and the bare residual as a diagnostic.

4.2 Composing with discovery: a 2×2

Because the residual is defined against the anchors, it can be built before or after the discovery loop. Writing A for the parent alone, B for the parent after discovery, C for the construction before discovery, and D for both, the quantity of interest is not $D - A$ but the interaction

$$\Delta_{\text{int}} = (D - C) - (B - A), \quad (7)$$

which is zero for an additive stack and positive only when the two ingredients do more together than apart. We report it for per-event power at a nominal false-positive rate $\alpha = 0.05$, because that metric is bounded below by α : a parent at $A \approx \alpha$ has *no* ability to flag an individual novel sample, so a positive interaction there is a threshold crossing rather than an improvement.

The interaction is real, and it is narrow. On DTD with a DINO backbone, the residual construction gives $A = 0.040$, $B = 0.068$, $C = 0.023$, $D = 0.253$. Discovery alone is worth +0.028; the construction alone is worth -0.018 — it actively hurts — and composed they reach $6\times$ nominal, an interaction of +0.203. The concatenation behaves the same way (+0.163). Neither ingredient does anything alone. This is the strongest evidence we have that discovery and representation construction are not independent improvements.

Two apparent replications are not replications. The same protocol on galaxy10 produces the two largest raw gains in the study, and the decomposition shows that neither is composition. With a DINO backbone the residual takes per-event power from 0.021 to 0.299, but $C = 0.281$ already: the construction does this *by itself*, and the interaction is +0.003. With a LeJEPa backbone the gain from 0.007 to 0.206 is almost entirely the discovery main effect (+0.190), with an interaction of +0.008 — and it occurs in a cell whose round-2 pool purity is exactly zero. That last case is worth stating plainly, because we had written it in advance as a falsifier: if the composition helps where purity is zero, then the mechanism we proposed for it — that discovery supplies a richer and *correct* anchor set, against which the residual becomes informative — cannot be what is happening there.

What we therefore claim. Across five cells, two constructions, and two variants, 12 of 20 interactions are positive, but most are small and the threshold crossings driven by interaction rather than by a main effect occur in a single cell with a single construction. We do not claim superadditivity as a general property. The claim the data supports is weaker and still useful: *discovery and the residual*

are substitutes more often than complements. Whichever of the two a cell is missing, one of them supplies it — galaxy10/LeJEPA needed discovery, galaxy10/DINO needed the residual, DTD/DINO needed both — and a practitioner who runs only one of them will, on most cells, capture most of the available gain. Every number in this subsection is a single seed; the interaction is a difference of differences and inherits the variance of both its terms, so we treat it as a finding to be replicated rather than a headline.

4.3 Frozen-space discovery

5 Experiments

How to read the tables. Four words recur and are worth fixing now. An *objective* (or arm) is one training loss from the design cube of Section 3 — SupCon, SupCon + SIGReg, and so on. A *backbone* (or base) is the pretrained trunk the objective is fine-tuned on top of: DINO, LeJEPA or VISReg. A *draw* is a choice of which class to hold out and treat as novel; re-drawing changes the problem, since some classes are far harder to find than others. A *seed* is a re-run of the same training with different random initialisation and batch order, holding the draw fixed. A *cell* is one (dataset, backbone) pair, inside which we compare objectives. Distinguishing draw from seed matters more than it may appear, and Section 5 measures both: they turn out to differ by almost an order of magnitude, which determines how much of a single-run number can be believed.

This campaign is reported while incomplete. Every table below states its own coverage, blanks are shown as “—” rather than dropped, and we mark which claims rest on a single seed. Datasets are ordered by *distance from pretraining* rather than alphabetically, because that ordering is what decides whether a result is evidence for discovery at all: galaxy10 was never seen by any backbone we use, DTD is task-novel, cars/flowers/aircraft are fine-grained but in-domain, and the CIFAR cells exist in two lineages — a from-scratch one that we report, and a hub-pretrained one whose encoder already saw the held-out class and which we therefore treat as an ablation only.

5.1 Objectives: the probe and the geometry come apart

Table 1 is the paper’s central measurement, on the lineage where it is safe to make it. Adding sliced-isotropy to a supervised contrastive objective is a tie on the linear probe and a 1.1-point loss on top-1 accuracy — and it more than doubles every discoverability column. A reader who ranked these eight objectives by probe accuracy would pick the one that scores 0.00 on per-event detection.

That is the dissociation stated at the start of the paper, measured on an encoder that was trained from random initialisation and never saw the class it is asked to find.

5.2 Stability: draws matter, seeds do not

Two variance questions decide how much of this campaign is quotable.

The first is the holdout *draw* — which class is held out. Table 2 reports five draws per cell. Pool purity is stable across draws for the best objective (sd 0.007–0.014) while the linear probe over the same draws swings by more than 0.1, so purity is the better-conditioned measurement of the two, not merely the more relevant one.

The second is the *seed*. We pre-registered this as a falsifier: if seed-to-seed spread were comparable to draw-to-draw spread, every single-seed number in this paper would need a caveat. Table 3 settles it on 144 cells — $\text{sd}_{\text{seed}} = 0.021$ against $\text{sd}_{\text{draw}} = 0.174$, a factor of 8. The caveat that survives is about which class you hold out, not which seed you ran. We attach it to every single-seed result rather than dropping it.

Table 2 also carries a result we did not anticipate. Only one objective holds its purity across *both* backbones. SupCon+SIGReg reaches 0.219 on DINO and 0.225 on LeJEPA, and its worst single draw (0.193) still exceeds every other objective’s mean. Supervised distance-NPLM matches it on LeJEPA (0.215) and loses a factor of four on DINO (0.051) — the same objective, the same data, the same code, differing only in the pretrained trunk. Base-independence is a stronger and more useful property than winning on average, and it is not visible in any single-backbone table.

objective	sup.	decodable			discoverable		
		probe	top-1	anchor	Eucl.	Maha.	per-event
SimCLR	n	0.799	0.292	0.242	0.443	0.385	0.01
VISReg/LeJEPa	n	0.797	0.239	0.190	0.499	0.382	0.03
NPLM (unsup.)	n	0.757	0.223	0.181	0.413	0.361	0.01
SupCon	y	0.940	0.608	0.545	0.233	0.279	0.00
SupCon+SIGReg (repulse)	y	0.823	0.360	0.346	0.340	0.320	0.00
NPLM-dist.+classwise SIGReg	y	0.898	0.507	0.463	0.525	0.411	0.03
SupCon+SIGReg ($\lambda=5$)	y	0.937	0.597	0.588	0.620	0.602	0.07
NPLM-dist. (sup.)	y	0.820	0.254	0.232	0.408	0.345	0.05

Table 1: Objectives on the leakage-free CIFAR-100 lineage (random init; the encoder never saw the held-out class). *Decodable* columns: linear-probe AUC for the held-out class (trained *with* its labels), top-1 accuracy on the seen classes, and nearest-anchor accuracy on the seen classes. *Discoverable* columns: novelty-ranking AUC by Euclidean and by tied-Mahalanobis distance, and per-event detection power at a 5% false-positive rate — the last is ≈ 0.05 for a space carrying no usable signal, so 0.00 means the space cannot flag an individual novel image at all. Adding SIGReg is a probe tie and a top-1 loss, and wins every discoverability column. *Coverage*: 8/8 objectives present; single holdout (class 4), one seed per cell, probe averaged over 3 probe restarts.

objective	DTD / DINO			DTD / LeJEPa		
	mean	sd	min	mean	sd	min
SimCLR	0.005	0.012	0.000	0.002	0.003	0.000
SimCLR+SIGReg	0.003	0.003	0.000	0.021	0.021	0.000
NPLM-bilinear	0.004	0.006	0.000	0.033	0.018	0.016
SupCon	0.028	0.014	0.011	0.073	0.086	0.011
SupCon+SIGReg ($\lambda=5$)	0.219	0.016	0.193	0.225	0.007	0.217
NPLM-dist. (sup.)	0.051	0.031	0.021	0.215	0.027	0.189

Table 2: Pool purity is stable across holdout draws, and only one objective is stable across backbones. Mean $\pm$ sd over draws, with the worst draw. SupCon+SIGReg holds its purity on both backbones (0.219 and 0.225); supervised distance-NPLM matches it on LeJEPa (0.215) and loses a factor of four on DINO (0.051). *Coverage*: DTD, 5 holdout draws per base, single holdout; 12/12 (arm, base) pairs present. Purity read from the run logs.

5.3 Constructions

Table 4 gives the 2×2 of Section 4 in full. The interaction is large in one cell, positive but small in most, and the two largest *totals* in the table are main effects rather than composition. We report all twenty rows rather than the favourable ones.

5.4 Where the label-free cut does not reach

Table 5 is a negative result about our own method. At a 1% novel rate on clean CIFAR-100 banks, the label-free rule of Section 3.3 engages on 3 of 20 spaces, and where it engages the resulting purity is 0.046 to 0.101 — an enrichment of 5 to 10 over the 1% base rate, but far short of what the same cut admits. The spaces where it works share a structure — an NPLM critic on a SIGReg marginal — so the rule should be read as conditional on the scorer, not as a general-purpose procedure. The ceiling algebra says a $20\times$ enrichment is available at this rate; the estimator, not the representation, is what fails to reach it.

6 When it does not work

A Detailed results

This appendix reports the campaign in more detail than the main text, and reports it *as it stands* rather than waiting for completion. All tables in this paper are generated from the archived result

round	strata	mean Π	sd_{seed}	sd_{draw}	ratio
r1	9	0.516	0.021	0.174	$8.3\times$
r2	9	0.517	0.020	0.175	$8.9\times$

Table 3: The variance that matters is the holdout draw, not the seed. One-way decomposition of frozen-anchor round-1 purity, stratified within (cell, objective) so the seed term is measured at fixed draw. Seed-to-seed spread is an eighth of draw-to-draw spread in both rounds, so the single-seed purities elsewhere in this paper carry a caveat about which class was held out and not about which seed was run. *Coverage:* galaxy10, 3 backbones $\times$ 3 objectives $\times$ 5 draws $\times$ 3 seeds (+1 archived) = 144 cells; variance stratified within (cell, objective). Round-level distribution summaries are not reported: the archive stores one summary for both rounds.

files by a single script, and each carries a machine-written coverage line stating what is present and what is missing. Blanks are shown as “—”.

galaxy10, and how much of the difficulty is the procedure. galaxy10 carries more weight than any other dataset here, because no backbone we use has seen it: DINO, LeJEPa and VISReg were all pretrained on natural images, and galaxy morphology is genuinely out of that distribution. It is also, for the same reason, the hardest cell in the study.

Table 6 reports it under both discovery procedures on the same holdout draws. Under natural discovery — a distance pool with the encoder fine-tuned — round-1 pool purity sits at 0.09–0.29 across every objective and backbone, with sds that are frequently as large as the means, so the objectives are not separated by this measurement at all. Under the frozen-anchor density-ratio pool, the same objectives on the same draws reach 0.39–0.58.

That is a factor of three to five from changing the procedure and nothing else. It is the sharpest evidence in the paper for a claim we make throughout: failures that present as properties of the representation are frequently properties of the pipeline wrapped around it. On galaxy10 specifically, the ordering of objectives under natural discovery is close to noise — the sds overlap almost completely — and the ordering under the frozen pool is both tighter and higher. We would not have found this by improving the encoder.

Note the coverage line: the frozen-pool arm of the study covers the three supervised objectives only, so nine of the eighteen cells in that half are absent rather than negative.

Constructions on the leakage-free trunks. Table 7 gives every construction on the from-scratch CIFAR-100 parents. The pattern behind the summary in Section 4 is visible directly: the bare residual usually costs probe accuracy, the concatenation recovers it, and the detection gain survives in both.

Multi-seed cells. Most cells in this campaign are single-seed, for cost reasons. Table 8 reports one of the exceptions, three seeds on the CIFAR-10 champions. Seed spread there is small — ± 0.001 on the best post-discovery probe and ± 0.002 on its Euclidean AUC — which is consistent with the much larger variance decomposition of Table 3. Note also that the discovery loop *reduces* the probe for supcon+np1m ($0.954 \rightarrow 0.930$) in all three seeds, which is the seeded version of a pattern we observe throughout: the loop repairs weak spaces and damages strong ones.

Known gaps. At the time of writing: the 2×2 of Table 4 is single-seed throughout, and its one large interaction rests on a single cell; the multi-holdout regime is run only on the label-rich datasets; and the frozen-space construction of Section 4 is quoted from re-verified runs after a BatchNorm fix invalidated the earlier archive, so its older numbers do not reproduce and are not used.

B Objective naming

cell	child loss	used as	A	B	C	D	$B-A$	$C-A$	Δ_{int}
cars / VISReg	SupCon	bare	0.223	0.171	0.073	0.066	-0.052	-0.150	+0.045
cars / VISReg	SupCon	concat	0.223	0.171	0.228	0.174	-0.052	+0.005	-0.002
cars / VISReg	NPLM	bare	0.223	0.171	0.256	0.284	-0.052	+0.033	+0.080
cars / VISReg	NPLM	concat	0.223	0.171	0.263	0.282	-0.052	+0.040	+0.070
DTD / DINO	SupCon	bare	0.040	0.068	0.022	0.253	+0.028	-0.018	+0.203
DTD / DINO	SupCon	concat	0.040	0.068	0.045	0.235	+0.028	+0.005	+0.163
DTD / DINO	NPLM	bare	0.040	0.068	0.113	0.087	+0.028	+0.073	-0.053
DTD / DINO	NPLM	concat	0.040	0.068	0.077	0.130	+0.028	+0.037	+0.025
galaxy10 / DINO	SupCon	bare	0.021	0.037	0.018	0.014	+0.015	-0.003	-0.020
galaxy10 / DINO	SupCon	concat	0.021	0.037	0.021	0.018	+0.015	-0.001	-0.018
galaxy10 / DINO	NPLM	bare	0.021	0.037	0.281	0.299	+0.015	+0.259	+0.003
galaxy10 / DINO	NPLM	concat	0.021	0.037	0.091	0.142	+0.015	+0.069	+0.036
galaxy10 / LeJEPa	SupCon	bare	0.007	0.196	0.001	0.002	+0.190	-0.005	-0.189
galaxy10 / LeJEPa	SupCon	concat	0.007	0.196	0.001	0.012	+0.190	-0.005	-0.179
galaxy10 / LeJEPa	NPLM	bare	0.007	0.196	0.009	0.021	+0.190	+0.002	-0.177
galaxy10 / LeJEPa	NPLM	concat	0.007	0.196	0.008	0.206	+0.190	+0.001	+0.008
galaxy10 / VISReg	SupCon	bare	0.311	0.419	0.243	0.457	+0.108	-0.068	+0.107
galaxy10 / VISReg	SupCon	concat	0.311	0.419	0.295	0.466	+0.108	-0.015	+0.063
galaxy10 / VISReg	NPLM	bare	0.311	0.419	0.333	0.471	+0.108	+0.022	+0.030
galaxy10 / VISReg	NPLM	concat	0.311	0.419	0.353	0.446	+0.108	+0.043	-0.015

Table 4: Discovery $\times$ construction, per-event power at $\alpha=0.05$. Each row is one 2×2 : A is the parent space alone, B the parent after the discovery loop, C the construction built on the parent *without* discovery, and D both together. The two main effects are $B - A$ (discovery) and $C - A$ (construction). The residual child is trained under the loss in the second column and then either used alone (*bare*) or concatenated with its parent (*concat*). The interaction $\Delta_{\text{int}} = (D - C) - (B - A)$ is zero if the two simply add, and positive only if they do more together than apart. Bold marks $\Delta_{\text{int}} > 0.10$. Note that the two largest *totals* in the table (both galaxy10) have near-zero interaction: they are main effects, not composition. *Coverage*: 5 cells $\times$ 2 constructions $\times$ 2 variants; 20 complete, 0 lacking a pre-discovery child. One seed per cell.

space	$\hat{b}/b$ (tv)	$\hat{b}/b$ (mass)	rule fires	q^*
nplm_cifar100	0.01	0.01	no	—
nplmcw_cifar100	1.04	0.49	yes	0.0277
nplmsd-res-nplm_(concat)_cifar100	0.15	0.08	no	—
nplmsd-res-nplm_(residual)_cifar100	0.07	0.04	no	—
nplmsd-res_(concat)_cifar100	0.17	0.14	no	—
nplmsd-res_(residual)_cifar100	0.37	0.28	no	—
nplmsd_cifar100	0.92	0.07	no	—
simclr_cifar100	0.09	0.09	no	—
ssig-res-nplm_(concat)_cifar100	0.18	0.04	no	—
ssig-res-nplm_(residual)_cifar100	0.43	0.36	no	—
ssig-res_(concat)_cifar100	0.04	0.03	no	—
ssig-res_(residual)_cifar100	0.01	0.01	no	—
ssig_cifar100	0.36	0.20	no	—
supcon-res-nplm_(concat)_cifar100	0.78	0.21	yes	0.0099
supcon-res-nplm_(residual)_cifar100	0.69	0.55	yes	0.0183
supcon-res_(concat)_cifar100	0.03	0.03	no	—
supcon-res_(residual)_cifar100	0.14	0.07	no	—
supcon_cifar100	0.00	0.01	no	—
supsig_cifar100	0.07	0.04	no	—
visreg_cifar100	0.47	0.40	no	—

Table 5: The label-free cut is objective-specific, not universal. Ratio of estimated to true novel base rate at $b=0.01$; the rule engages only where an NPLM critic sits on a SIGReg marginal, and even then the resulting purity is only 0.046–0.101 (against a base rate of 0.01). Reported as a limit of the rule, not of the representations. *Coverage*: 20 spaces, clean CIFAR-100 banks, $b_{\text{true}}=0.01$; the rule engages on 3, and the best estimator is within $5\times$ on 8.

objective	natural discovery			frozen-anchor pool		
	DINO	LeJEPa	VISReg	DINO	LeJEPa	VISReg
SimCLR	0.104±0.089	0.112±0.100	0.093±0.063	—	—	—
SimCLR+SIGReg	0.100±0.088	0.123±0.168	0.110±0.077	—	—	—
NPLM-bilinear	0.104±0.080	0.102±0.051	0.117±0.092	—	—	—
SupCon	0.103±0.138	0.132±0.085	0.290±0.067	0.465±0.158	0.575±0.185	0.574±0.180
SupCon+SIGReg ($\lambda=5$)	0.092±0.074	0.182±0.200	0.146±0.042	0.459±0.143	0.516±0.186	0.554±0.169
NPLM-dist. (sup.)	0.131±0.038	0.190±0.137	0.262±0.106	0.393±0.147	0.569±0.179	0.543±0.185

Table 6: galaxy10, the dataset no backbone has seen, under two discovery procedures. Round-1 pool purity, mean $\pm$ sd. Left: natural discovery (distance pool, encoder fine-tuned). Right: the frozen-anchor density-ratio pool. Same dataset, same holdout draws, a three- to five-fold difference in purity — the procedure, not the representation, is what moves it. *Coverage*: galaxy10, 3 backbones; natural discovery over 5 draws (18/18 cells present), frozen-anchor pool over 5 draws $\times$ 3 seeds (9/18 present — that arm of the study covers the three supervised objectives only).

construction	probe	Δ	Eucl. AUC	Δ	per-event
<i>parent: SupCon</i>					
(parent)	0.940	–	0.233	–	0.00
residual child (SupCon), bare	0.869	-0.071	0.607	+0.374	0.05
residual child (SupCon), concat	0.942	+0.002	0.521	+0.288	0.03
residual child (NPLM), bare	0.882	-0.058	0.291	+0.057	0.00
residual child (NPLM), concat	0.947	+0.007	0.227	-0.006	0.00
<i>parent: SupCon+SIGReg ($\lambda=5$)</i>					
(parent)	0.937	–	0.620	–	0.07
residual child (SupCon), bare	0.860	-0.077	0.604	-0.017	0.08
residual child (SupCon), concat	0.949	+0.012	0.648	+0.028	0.08
residual child (NPLM), bare	0.900	-0.037	0.464	-0.156	0.01
residual child (NPLM), concat	0.904	-0.033	0.575	-0.046	0.04
<i>parent: NPLM-dist. (sup.)</i>					
(parent)	0.820	–	0.408	–	0.05
residual child (SupCon), bare	0.834	+0.014	0.524	+0.116	0.01
residual child (SupCon), concat	0.867	+0.047	0.521	+0.114	0.01
residual child (NPLM), bare	0.827	+0.006	0.527	+0.119	0.06
residual child (NPLM), concat	0.849	+0.029	0.516	+0.108	0.04

Table 7: All constructions on the leakage-free CIFAR-100 trunks. Δ columns are against that parent. The bare residual generally loses probe accuracy and the concatenation recovers it; the detection gain is large and nearly free. *Coverage:* 3 parents $\times$ 4 constructions on the from-scratch CIFAR-100 trunks; one seed per cell.

space	probe (pre)	probe (post)	Eucl. AUC	Maha.	seeds
nplmsup+nplm	0.888 $\pm$ 0.010	0.917 $\pm$ 0.004	0.686 $\pm$ 0.016	0.752 $\pm$ 0.025	3
nplmsup- ζ res	0.871 $\pm$ 0.015	0.879 $\pm$ 0.017	0.448 $\pm$ 0.015	0.621 $\pm$ 0.045	3
sup+nplm	0.926 $\pm$ 0.012	0.979 $\pm$ 0.001	0.794 $\pm$ 0.009	0.729 $\pm$ 0.044	3
sup- ζ res-nplm	0.930 $\pm$ 0.012	0.983 $\pm$ 0.001	0.834 $\pm$ 0.002	0.759 $\pm$ 0.012	3
supcon+nplm	0.954 $\pm$ 0.002	0.930 $\pm$ 0.009	0.766 $\pm$ 0.005	0.723 $\pm$ 0.016	3

Table 8: Multi-seed CIFAR-10 champions. Mean $\pm$ sd over seeds. Seed spread here is small, consistent with Table 3. *Coverage:* CIFAR-10, 5 spaces, up to 3 seeds; one of the few multi-seed cells in the campaign.